AI Auto Insurance Tool — Agent Requirements

REQ-001: Policy Coverage Parser

Goal: Extract key auto insurance coverage details from uploaded policy documents.

Agent task:
Build a parser that identifies liability, collision, comprehensive, uninsured/underinsured motorist, rental reimbursement, medical payments, deductibles, limits, exclusions, endorsements, and effective dates.

Acceptance criteria:

- User can upload or paste a policy.
- System extracts coverage sections into structured JSON.
- Each extracted field links back to the original policy text.
- Missing or unclear fields are flagged instead of guessed.

REQ-002: Coverage Explanation Engine

Goal: Explain what the user is covered for in plain English.

Agent task:
Create an explanation layer that converts policy language into user-friendly summaries.

Acceptance criteria:

- Each coverage type has a plain-English explanation.
- The system explains what is covered, what is not covered, and what may require review.
- The explanation avoids legal certainty unless supported by the policy.
- Output includes confidence level: high, medium, or low.

REQ-003: Scenario-Based Coverage Check

Goal: Let users ask, “Am I covered if this happens?”

Agent task:
Build a scenario evaluator that compares a user’s accident scenario against policy coverage.

Acceptance criteria:

- User can describe a real-world situation.

- System identifies relevant coverages and exclusions.
- System returns likely covered, likely not covered, or needs human review.
- System cites supporting policy sections.

REQ-004: Temporal Coverage Timeline

Goal: Determine whether coverage existed at a specific date and time.

Agent task:
Build a policy timeline engine that tracks effective dates, expiration dates, cancellations, reinstatements, endorsements, and gaps.

Acceptance criteria:

- User can ask, “Was I covered on this date?”
- System checks the policy timeline.
- System flags lapses, overlaps, and unclear coverage periods.
- System handles bridge coverage between old and new policies.

REQ-005: Bridge Policy / Coverage Continuity Logic

Goal: Detect and explain coverage transitions between policies.

Agent task:
Build bridge logic that compares prior policy, new policy, cancellation date, start date, and claim date.

Acceptance criteria:

- System identifies whether there is a gap between policies.
- System identifies whether two policies overlap.
- System explains which policy may respond to a claim.
- System flags unclear transitions for human review.

REQ-006: Claim Intake Assistant

Goal: Improve First Notice of Loss intake.

Agent task:
Build an adaptive claim intake assistant that asks only relevant follow-up questions based on the accident type.

Acceptance criteria:

- User can start a claim from a conversational form.
- System collects date, time, location, driver, vehicle, damage, injuries, police report, photos, and third-party details.
- Questions change based on the user's answers.
- System produces a structured claim summary.

REQ-007: Claim Coverage Decision Draft

Goal: Generate an explainable preliminary claim coverage opinion.

Agent task:
Create a claim review module that maps the claim facts to the policy.

Acceptance criteria:

- System identifies relevant coverages.
- System identifies possible exclusions.
- System explains likely claim outcome.
- System separates facts, assumptions, policy language, and reasoning.
- System never presents the AI result as a final legal or carrier decision.

REQ-008: Out-of-Pocket Cost Estimator

Goal: Show the user what they may actually pay.

Agent task:
Build an estimator for deductible, limits, rental reimbursement, uncovered costs, and possible user responsibility.

Acceptance criteria:

- User can enter estimated repair cost.
- System applies deductible and policy limits.
- System estimates likely user out-of-pocket cost.
- System clearly explains assumptions.

REQ-009: Coverage Gap Detector

Goal: Identify risky missing or weak coverages.

Agent task:
Build a detector that reviews the policy and flags gaps such as missing rental coverage, low liability limits, missing UM/UIM, high deductibles, or excluded usage.

Acceptance criteria:

- System lists detected gaps.
- Each gap explains why it matters.
- Each gap includes potential real-world consequence.
- System avoids upselling language unless explicitly configured.

REQ-010: Edge Case Classifier

Goal: Handle messy claim situations carriers and tools often miss.

Agent task:
Create classifiers for non-listed drivers, rideshare/delivery use, out-of-state accidents, borrowed vehicles, business use, permissive use, excluded drivers, and multi-car accidents.

Acceptance criteria:

- System detects when a scenario contains an edge case.
- System routes the claim to enhanced review.
- System explains why the case is complex.
- System identifies which policy sections need review.

REQ-011: Policy Evidence Mapper

Goal: Make every AI answer traceable.

Agent task:
Build a citation layer that maps every important answer to policy text.

Acceptance criteria:

- Every coverage conclusion includes source text.
- Citations include page, section, or clause when available.
- Unsupported claims are marked as assumptions.
- Users can click or view the original language.

REQ-012: Confidence and Escalation System

Goal: Prevent the AI from over-answering uncertain insurance questions.

Agent task:
Build confidence scoring and escalation rules.

Acceptance criteria:

- Each answer receives a confidence rating.
- Low-confidence answers trigger “human review recommended.”
- Conflicting policy language is flagged.
- Missing documents or missing facts are clearly listed.

REQ-013: State-Specific Rules Layer

Goal: Account for state-level insurance differences.

Agent task:
Create a rules framework that can apply state-specific auto insurance rules.

Acceptance criteria:

- Policy state is detected or entered by user.
- System can attach jurisdiction-specific warnings.
- State rules are separated from AI-generated reasoning.
- Rules are versioned and auditable.

REQ-014: Adjuster / Agent Copilot View

Goal: Give professionals a structured review interface.

Agent task:
Build a professional-facing dashboard that summarizes claim facts, coverage issues, risk flags, missing information, and AI reasoning.

Acceptance criteria:

- Dashboard shows claim summary.
- Dashboard shows policy sections involved.
- Dashboard shows possible claim outcome.
- Dashboard shows unresolved questions.
- Human reviewer can approve, reject, or override AI suggestions.

REQ-015: Audit Trail

Goal: Record how every answer was produced.

Agent task:
Build an audit log for policy uploads, extracted data, user questions, AI responses, source citations, confidence levels, and human overrides.

Acceptance criteria:

- Every AI response is logged.
- Every source document version is tracked.
- Human overrides are stored.
- Logs can be exported for compliance review.

REQ-016: Feedback Loop

Goal: Improve the system over time using real outcomes.

Agent task:
Build a feedback system where users or reviewers can mark AI answers as correct, incorrect, incomplete, or unclear.

Acceptance criteria:

- User can rate an answer.
- Adjuster can override a recommendation.
- Feedback is stored with the original policy, scenario, and AI answer.
- Feedback can be reviewed later for model improvement.

REQ-017: Safe Failure Mode

Goal: Make the AI honest when it cannot answer.

Agent task:
Build a safe fallback response system.

Acceptance criteria:

- System says when it cannot determine coverage.
- System lists exactly what information is missing.
- System recommends next action.

- System does not fabricate policy terms or claim outcomes.

REQ-018: Policy Comparison Tool

Goal: Compare two auto policies side by side.

Agent task:
Build a comparison module for old policy vs new policy or current policy vs quoted policy.

Acceptance criteria:

- System compares limits, deductibles, exclusions, endorsements, premiums, and effective dates.
- System highlights coverage improvements and reductions.
- System identifies possible bridge gaps between policies.
- System produces a plain-English summary.

REQ-019: User Question Answering

Goal: Let users ask natural-language questions about their auto policy.

Agent task:
Build a question-answering interface grounded only in the uploaded policy and approved rules.

Acceptance criteria:

- User can ask free-text questions.
- System answers using policy-grounded reasoning.
- System cites the relevant policy text.
- System refuses or escalates when the policy does not support an answer.

REQ-020: MVP Priority Order

Build in this order:

1. Policy Coverage Parser
2. Policy Evidence Mapper
3. Coverage Explanation Engine
4. Temporal Coverage Timeline
5. Bridge Coverage Logic
6. Scenario-Based Coverage Check
7. Claim Intake Assistant
8. Claim Coverage Decision Draft

9. Confidence and Escalation System
10. Audit Trail